

Machine Learning for Time Series

Lecture 1: Pattern Recognition and Detection

Laurent Oudre

laurent.oudre@ens-paris-saclay.fr

Master MVA

2026-2027

Contents

1. Problem statement

2. Comparing time series

2.1 Euclidean distance

2.2 Normalized Euclidean distance

2.3 Dynamic Time Warping

3. Detecting patterns in time series

3.1 Euclidean distance

3.2 DTW

4. Learning patterns from time series

4.1 Distance-based pattern extraction

4.2 Dictionary-based pattern extraction


5. Conclusion

Contents


1. Problem statement
2. Comparing time series
3. Detecting patterns in time series
4. Learning patterns from time series
5. Conclusion

Motivation



Power usage consumer profiling. Source / US National Institute of Standards and Technology




Motivation



Problem 1: Pattern Detection





Dictionary of patterns




Input time series

Problem 1: Pattern Detection



Dictionary of patterns



Annotated time series


Problem 1: Pattern Detection

Pattern Detection



Given a dictionary of patterns, retrieve these patterns in an input time series

- ▶ The patterns/the time series can be **multivariate**
- ▶ The patterns may have **different lengths**
- ▶ The patterns can be **annotated**, i.e., be linked to a specific phenomenon of interest: in this context, pattern recognition will provide an automated annotation of the input time series

Problem 2: Pattern Extraction



Input time series



Extracted patterns

Problem 2: Pattern Extraction

Pattern Extraction

Given an input time series (or a set of time series), learn a dictionary of patterns

- ▶ A pattern is a **shape** that appears repetitively in the time series (though a somewhat blurry notion)
- ▶ All patterns are supposed to have the **same length** (for the sake of simplicity)
- ▶ The extracted patterns can be used to **characterize** the time series, or studied individually

Contents

1. Problem statement

2. Comparing time series

2.1 Euclidean distance

2.2 Normalized Euclidean distance

2.3 Dynamic Time Warping


3. Detecting patterns in time series

4. Learning patterns from time series


5. Conclusion

Invariances and similarity measures

- ▶ To determine whether a pattern repeats, we need a way to **compare time series** or subsequences of time series. To do this, we will rely on **similarity measures**.
- ▶ Occurrences of the same pattern are **never identical**. They may differ by:
 - ▶ **Offset** : $y = x + \beta \mathbf{1}$
 - ▶ **Amplitude scaling** : $y = \alpha x, \alpha > 0$
 - ▶ **Linear trend** : $y[n] = x[n] + an + b$
 - ▶ **Time shift** : $y[n] = x[n - \tau]$
 - ▶ **Uniform scaling** : longer or shorter occurrence
 - ▶ **Local time warping** : non-uniform stretching
 - ▶ **Additive noise and outliers**
- ▶ A similarity measure is useful to the extent that it is **blind** to the deformations deemed irrelevant, and **sensitive** to everything else.
- ▶ There is no universally correct choice: this is an **application-level decision**.



Euclidean distance




$$d_{EUC}(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{n=1}^N (x[n] - y[n])^2}$$

- ▶ **Sensitive** to time shifts, amplitude changes, offsets and dilatation/contraction
- ▶ Necessity to have a **perfect match** between the timelines
- ▶ Sensitive to **outliers** but approximately OK with low AWGN


What influences Euclidean distance?

Baseline case



$$d_{EUC} = 3.5$$


Offset



$$d_{EUC} = 9.6$$


What influences Euclidean distance?

Amplitude scaling (70 %)



$$d_{EUC} = 12.8$$


Time shift (1.6 %)



$$d_{EUC} = 5.3$$


What influences Euclidean distance?

Additive white Gaussian noise



$$d_{EUC} = 6.9$$

Outliers



$$d_{EUC} = 10.3$$

Normalized Euclidean distance

In order to improve the robustness to changes in amplitude/offset, several authors recommend using a **normalized Euclidean distance** :

$$d_{nEUC}(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{n=1}^N (\tilde{x}[n] - \tilde{y}[n])^2}$$


where

$$\tilde{x}[n] = \frac{x[n] - \mu_x}{\sigma_x} \quad \text{z-score normalization}$$

- ▶ **Still sensitive** to time shifts, dilatation/contraction, outliers and misaligned timelines
- ▶ Normalization may **increase the sensitivity** with respect to additive noise


What influences normalized Euclidean distance?

Baseline case



$$d_{nEUC} = 0.8$$


Offset



$$d_{nEUC} = 0.8$$


What influences normalized Euclidean distance?

Amplitude scaling (70 %)



$$d_{nEUC} = 0.8$$


Time shift (1.6 %)



$$d_{nEUC} = 1.9$$


What influences normalized Euclidean distance?

Additive white Gaussian noise



$$d_{nEUC} = 1.9$$

Outliers



$$d_{nEUC} = 2.7$$

Impact of the normalization

- ▶ Expanding the z-normalized distance gives

$$d_{nEUC}^2(\mathbf{x}, \mathbf{y}) = 2N(1 - \rho) \quad \text{with} \quad \rho = \frac{\sum_n x[n]y[n] - N\mu_x\mu_y}{N\sigma_x\sigma_y}$$

where ρ is the Pearson correlation coefficient

- ▶ **Ranking by z-normalized distance = ranking by Pearson correlation**
- ▶ The distance depends on the data only through the **inner product** and first/second-order moments
- ▶ Dividing by σ_x **amplifies noise** when the subsequence is nearly flat: a constant segment with small noise becomes pure amplified noise, which then matches arbitrary shapes

Extension to linear trends

- ▶ z-normalization removes a constant offset, **but not a linear drift**
- ▶ It is possible to use more refined normalization process, for instance to remove linear trend: this can be done by removing the projection on $\mathcal{A} = \text{span} \{ \mathbf{1}, (1, 2, \dots, N)^\top \}$
- ▶ LT-normalization can be defined as:

$$\mathbf{x}^{LT} = \frac{(I - \Pi_{\mathcal{A}})\mathbf{x}}{\|(I - \Pi_{\mathcal{A}})\mathbf{x}\|_2} \quad d_{LT}(\mathbf{x}, \mathbf{y}) = \|\mathbf{x}^{LT} - \mathbf{y}^{LT}\|_2$$

- ▶ It can be computed as easily as z-normalization, and z-normalization can be seen as a particular case of LT-normalization
- ▶ The same principle can be used on larger invariance groups [[Germain et al., 2024](#)] : vast literature on the notion of invariances/deformations in time series


Open questions

- ▶ How do we decrease the sensitivity with respect to timelines (contraction/dilatation, time shifts...)?
- ▶ How can we compare two time series of different lengths?
- ▶ How do we extend the notion of Euclidean distance to nonlinear timeline modifications?


New similarity measure : Dynamic Time Warping (DTW) [Berndt et al., 1994]

Differences between Euclidean and DTW

Euclidean distance




DTW



Sample $x[n]$ is associated to sample $y[n]$ Sample $x[i_k]$ is associated to sample $y[j_k]$

Notion of path



- ▶ DTW computes a correspondence between the elements of $\mathbf{x}$ and those of $\mathbf{y}$.
- ▶ Mapping function called **path** :

$$P = ((i_1, j_1), \dots, (i_{K_P}, j_{K_P})) \in (\mathbb{N} \times \mathbb{N})^{K_P}, K_P \in \mathbb{N}$$

$$(i_k, j_k) \in P \Leftrightarrow y[j_k] \text{ is matched with } x[i_k]$$

Choice of the optimal path

- ▶ The path P is evaluated through a **cost function**

$$w(P) = \sum_{k=1}^{K_P} (x[i_k] - y[j_k])^2$$

- ▶ The final DTW distance is computed as the **minimal value** for the cost function

$$d_{DTW}(\mathbf{x}, \mathbf{y}) = \sqrt{\min_{P \in \mathcal{P}} w(P)}$$

- ▶ Note that if $M = N$, $K_P = N$ and $i_k = j_k = k$, DTW is exactly equal to the Euclidean distance

Set of acceptable paths $\mathcal{P}$

How do we choose the set of acceptable paths $\mathcal{P}$? Given two time series $\mathbf{x}$ and $\mathbf{y}$ of lengths M and N , the path P should be:

- ▶ **Continuous** $i_k - i_{k-1} \leq 1$ and $j_k - j_{k-1} \leq 1$

Imagining the plane as a chessboard, the path can only go to the positions a king could move to;

- ▶ **Monotonic** $i_{k-1} \leq i_k$ and $j_{k-1} \leq j_k$


The path can only go up and right, and diagonally up and right;

- ▶ **Bounded** $(i_1, j_1) = (1, 1)$ and $(i_{K_P}, j_{K_P}) = (M, N)$

The path starts/ends by matching together the first/last elements of both signals

$\mathcal{P}$: set of paths respecting those three restrictions

Set of acceptable paths $\mathcal{P}$



Limited set of indexes to get to (i_k, j_k)

$$(i_{k-1}, j_{k-1}) = \begin{cases} (i_k - 1, j_k) \\ \text{or } (i_k, j_k - 1) \\ \text{or } (i_k - 1, j_k - 1). \end{cases}$$

A recursive problem

- ▶ Let $C(i, j)$ be the minimum cost for the path starting at $(1, 1)$, arriving at (i, j) , and belonging to $\mathcal{P}$
- ▶ It goes that

$$\text{DTW}(\mathbf{x}, \mathbf{y}) = \sqrt{C(M, N)}$$

- ▶ So how can we compute $C(M, N)$? By using a **dynamic programming** approach
- ▶ By using the properties of the acceptable paths, we can compute this quantity **recursively**

DTW algorithm

► Initialization:

$$C(1, j) = \sum_{\ell=1}^j (x[1] - y[\ell])^2 \quad C(i, 1) = \sum_{\ell=1}^i (x[\ell] - y[1])^2$$

► How to compute $C(2, 2)$? There are three ways to reach $(2, 2)$

- From $(1, 1)$: associated cost $C(1, 1) + (x[2] - y[2])^2$
- From $(1, 2)$: associated cost $C(1, 2) + (x[2] - y[2])^2$
- From $(2, 1)$: associated cost $C(2, 1) + (x[2] - y[2])^2$

► If we only keep the minimum cost, we have

$$C(2, 2) = (x[2] - y[2])^2 + \min \{C(1, 1), C(1, 2), C(2, 1)\}$$

► It is therefore possible to construct the matrix $C(i, j)$ iteratively with increasing i and j

DTW algorithm

Algorithm 1: Dynamic Time Warping

Inputs : Time series $\mathbf{x} \in \mathbb{R}^M$ and $\mathbf{y} \in \mathbb{R}^N$

Output: DTW $d_{DTW}(\mathbf{x}, \mathbf{y})$

$\mathbf{D} = \mathbf{0}_{M \times N}$;

$\mathbf{C} = \mathbf{0}_{M \times N}$;

for $1 \leq i \leq M, 1 \leq j \leq N$ **do**

$D(i, j) = (x[i] - y[j])^2$;

end

for $1 \leq i \leq M, 1 \leq j \leq N$ **do**

$C(1, j) = \sum_{\ell=1}^j D(1, \ell)$;

$C(i, 1) = \sum_{\ell=1}^i D(\ell, 1)$;

end


for $2 \leq i \leq M, 2 \leq j \leq N$ **do**

$C(i, j) = D(i, j) + \min \{C(i-1, j-1), C(i-1, j), C(i, j-1)\}$;



end

$d_{DTW}(\mathbf{x}, \mathbf{y}) = \sqrt{C(M, N)}$;


Step 1 : Distance matrix




Step 2 : Cumulative distance matrix



Step 2 : Cumulative distance matrix



Fast computation and variants



- ▶ If the timelines of $\mathbf{x}$ and $\mathbf{y}$ are assumed to be approximately similar, it is not necessary (nor advisable) to compute all values of $C(i, j)$
- ▶ Constrained variants can be implemented, by for instance only computing $C(i, j)$ if


$$|i - j| < \lambda \text{ (Sakoe-Chiba band)}$$

[Sakoe et al., 1978]

- ▶ Other truncations exist such as the **Itakura Parallelogram** [Itakura, 1975] and depend on the assumptions made on the type of nonlinear timeline variations


Fast computation and variants

Classic DTW



9 ms

Constrained DTW



2 ms

What influences DTW?

- ▶ By definition, DTW is **robust to timeline changes**
- ▶ However, it still requires the first and last samples to be aligned, which can be a limitation in case of time shifts. Some variants exist to avoid this problem.
- ▶ DTW is **sensitive to scale** (can be solved by normalizing the time series beforehand), noise and outliers

Pros and cons of DTW

- ▶ d_{DTW} is symmetric and equal to zero on identical inputs, but it **violates the triangle inequality**
- ▶ **Not a metric :**
 - ▶ Forbids metric indexing structures, complicates the definition of barycenters, degrades k -means-type clustering
 - ▶ Other elastic distances based on “edit transformations” are real distances : **TWED (Time Warp Edit Distance)**, **MSM (Move-Split-Merge)**... (cf mini-projects)
- ▶ **Not differentiable :**
 - ▶ Not usable as a training loss for a neural network
 - ▶ Soft variants (with soft-min instead of min) exist : **soft-DTW**, **DILATE**, **soft-MSM**... (cf mini-projects)


Summary

	Euclidean	Normalized Euclidean	DTW	Normalized DTW
Offset	✗	✓	✗	✓
Amplitude shift	✗	✓	✗	✓
Time shift	✗	✗	≈	≈
Dilatation or contraction	✗	✗	✓	✓
Additive noise	✗	✗	✗	✗
Outlier	✗	✗	✗	✗

Contents

- 1. Problem statement
- 2. Comparing time series
- 3. Detecting patterns in time series
 - 3.1 Euclidean distance
 - 3.2 DTW
- 4. Learning patterns from time series
- 5. Conclusion

Problem statement



- In order to find a known pattern $\mathbf{p}$ of length N_p in a time series $\mathbf{x}$ of length N , we need to compute all distances

$$d[n] = d(\mathbf{p}, x[n : n + N_p - 1])$$

for $1 \leq n \leq N - N_p + 1$ and where $x[n : n + N_p - 1]$ is the sequence of $\mathbf{x}$ starting at n and of length N_p

- Intuitively, this is a costly operation since the distance needs to be computed approximately N times

Can we compute this quantity with good complexity?

Fast computation for Euclidean distance

$$d[n] = d(\mathbf{p}, x[n : n + N_p - 1])$$

- ▶ The sequence $d[n]$ is referred to in the literature as the **distance profile**
- ▶ A solution has been found for standard and normalized Euclidean distance. The whole computation is in $\mathcal{O}(N \log N)$ and does not depend on the pattern length N_p [Mueen et al., 2015]
- ▶ This solution is based on the properties of **FFT (Fast Fourier Transform)** that is a fundamental algorithm for computing the Discrete Fourier Transform of a time series in $\mathcal{O}(N \log N)$

Reformulation of the Euclidean distance

- Given a pattern $\mathbf{p}$ of length N_p and a time series $\mathbf{x}$ of length $N > N_p$, we can write

$$d_{EUC}(\mathbf{p}, \mathbf{x}[n : n + N_p - 1]) = \sqrt{\sum_{i=n}^{n+N_p-1} x[i]^2 + \sum_{i=1}^{N_p} p[i]^2 - 2 \sum_{i=1}^{N_p} x[n+i-1]p[i]}$$

$$d_{nEUC}(\mathbf{p}, \mathbf{x}[n : n + N_p - 1]) = \sqrt{2N_p \left(1 - \frac{\sum_{i=1}^{N_p} x[n+i-1]p[i] - N_p \mu_x[n] \mu_p}{N_p \sigma_x[n] \sigma_p} \right)}$$

where

$$\mu_x[n] = \frac{1}{N_p} \sum_{i=n}^{n+N_p-1} x[i], \quad \sigma_x[n] = \sqrt{\frac{1}{N_p} \sum_{i=n}^{n+N_p-1} (x[i] - \mu_x[n])^2}$$

are sliding mean/standard deviations

- All statistics related to $\mathbf{p}$ such as $\sum_{i=1}^{N_p} p[i]^2$, μ_p or σ_p are easy to compute, but those of $\mathbf{x}[n : n + N_p - 1]$ change for each index n

Reformulation of the Euclidean distance

- ▶ The quantities involving $x[n : n + N_p - 1]$ to be computed are

$$\begin{aligned}s_x[n] &= \sum_{i=n}^{n+N_p-1} x[i]^2 \\ \mu_x[n] &= \frac{1}{N_p} \sum_{i=n}^{n+N_p-1} x[i] \\ \sigma_x[n] &= \sqrt{\frac{1}{N_p} \sum_{i=n}^{n+N_p-1} (x[i] - \mu_x[n])^2}\end{aligned}$$

- ▶ In fact, by pre-computing (in linear time) and storing the cumulative terms

$$\begin{aligned}c_x[n] &= \sum_{i=1}^n x[i] \\ c_x^{(2)}[n] &= \sum_{i=1}^n x[i]^2\end{aligned}$$

all quantities can be computed in a few operations

Reformulation of the Euclidean distance

- The quantities involving $x[n : n + N_p - 1]$ to be computed are

$$\begin{array}{l|l}
 s_x[n] = \sum_{i=n}^{n+N_p-1} x[i]^2 & s_x[n] = c_x^{(2)}[n + N_p - 1] - c_x^{(2)}[n - 1] \\
 \mu_x[n] = \frac{1}{N_p} \sum_{i=n}^{n+N_p-1} x[i] & \mu_x[n] = \frac{1}{N_p} (c_x[n + N_p - 1] - c_x[n - 1]) \\
 \sigma_x[n] = \sqrt{\frac{1}{N_p} \sum_{i=n}^{n+N_p-1} (x[i] - \mu_x[n])^2} & \sigma_x[n] = \sqrt{\frac{1}{N_p} s_x[n] - \mu_x[n]^2}
 \end{array}$$

- In fact, by pre-computing (in linear time) and storing the cumulative terms

$$\begin{aligned}
 c_x[n] &= \sum_{i=1}^n x[i] \\
 c_x^{(2)}[n] &= \sum_{i=1}^n x[i]^2
 \end{aligned}$$

all quantities can be computed in a few operations

Fast computation of the dot product

- ▶ The only difficulty is to compute the dot product

$$r[n] = \sum_{i=1}^{N_p} x[n+i-1]p[i]$$

- ▶ We use here two important results from signal processing :
 - ▶ Sliding dot product corresponds to the **discrete convolution** between the time series $\mathbf{x}$ and the time series $\mathbf{p}^\top$ where $\mathbf{p}^\top$ is the time-reversed time series $\mathbf{p}$
 - ▶ Convolution in the time domain corresponds to a **term-by-term product** in the frequency domain (up to some zero padding). It is always cheaper to compute a convolution product by computing the inverse FFT of the term-by-term product of the FFTs

$$\mathbf{x} * \mathbf{y} = \text{iFFT} \{ \text{FFT} \{ \mathbf{x} \} \times \text{FFT} \{ \mathbf{y} \} \}$$

Fast computation of the dot product

1. Compute the zero-padded time-reversed pattern

$$\mathbf{p}^{\leftarrow} = p[N_p]p[N_p - 1] \dots p[1] \underbrace{0 \dots 0}_{N - N_p}$$

2. The dot product can be computed as

$$\mathbf{r}[1 : N - N_p + 1] = \text{iFFT} \left\{ \text{FFT} \{ \mathbf{x} \} \times \text{FFT} \{ \mathbf{p}^{\leftarrow} \} \right\} [N_p : N]$$

Thanks to the fast implementations of FFT, this step is only $\mathcal{O}(N \log N)$

Final computation

- ▶ Both computations of sliding Euclidean distance and normalized Euclidean distance can be computed in $\mathcal{O}(N \log N)$
- ▶ Normalized Euclidean distance is **not more computationally demanding** than standard Euclidean distance: the only difference lies in the computation of $\mu_x[n]$ which is linear
- ▶ With $N_p = 10^2$ and $N = 10^8$, computation time is only 16 seconds on standard computers

Normalized Euclidean distance profile

Algorithm 2: Normalized Euclidean distance profile

Inputs : Time series $\mathbf{x} \in \mathbb{R}^N$ and pattern $\mathbf{p} \in \mathbb{R}^{N_p}$

Output: Normalized Euclidean distance profile $\mathbf{d} \in \mathbb{R}^{N-N_p+1}$

Normalize $\mathbf{p}$ (zero-mean, unit variance) and compute $\mathbf{p}^\dagger = p[N_p]p[N_p - 1] \dots p[1] \underbrace{0 \dots 0}_{N-N_p}$;

$\mathbf{c}_1 = \mathbf{0}_{N+1}$; $\mathbf{c}_2 = \mathbf{0}_{N+1}$;

for $2 \leq i \leq N + 1$ **do**

$c_1(i) = c_1(i - 1) + x[i - 1]$;
 $c_2(i) = c_2(i - 1) + x[i - 1]^2$;

end

$\mathbf{r} = \text{iFFT} \{ \text{FFT} \{ \mathbf{x} \} \times \text{FFT} \{ \mathbf{p}^\dagger \} \}$;




$\mathbf{s} = c_2(N_p + 1 : N + 1) - c_2(1 : N - N_p + 1)$;

$\boldsymbol{\mu} = (c_1(N_p + 1 : N + 1) - c_1(1 : N - N_p + 1)) / N_p$;

$\boldsymbol{\sigma} = \sqrt{\mathbf{s} / N_p - \boldsymbol{\mu}^2}$;

$\mathbf{d} = \sqrt{2(N_p - r(N_p : N) / \boldsymbol{\sigma})}$;

Results



Patterns can easily be retrieved in the signal by searching for peaks with small distance values

Fast computation for DTW

- ▶ The problem is more complex for DTW and will be reformulated as follows :

$$\text{Find } d^* = \min_{1 \leq n \leq N - N_p + 1} d_{DTW}(\mathbf{p}, x[n : n + N_p - 1])$$

- ▶ Instead of computing all DTW distances, we are searching for the **minimal DTW** between the pattern and all subsequences
- ▶ Similarly to Euclidean distance, a trick with pre-computed cumulative sums can be used to compute the DTW between renormalized signals (zero mean, unit variance), in order to be robust to amplitude and offsets.
- ▶ A very fast solution exists, that is based on **several acceleration tricks** [Rakthanmanon et al., 2012], one of them being the existence of a **lower bound**

Finding the minimal DTW distance

The problem here is different than for Euclidean distance: we are only interested in the **minimal distance**:

- ▶ **First trick** : use the lower bound for DTW. If we know that $d_{DTW}(\mathbf{x}, \mathbf{y}) \geq LB(\mathbf{x}, \mathbf{y})$ and that LB is cheap to compute, then when LB is larger than the current minimal distance, it is not worth computing the DTW.
- ▶ **Second trick** : even when computing a lower bound LB (which is based on the Euclidean distance : see next slides) or the DTW we can stop when the sum on the first samples exceeds the current minimal distance
- ▶ **Other tricks** : other tricks (such as using a pruned computation for the normalization, re-ordering the sequences or using cascade lower bounds) can help to avoid more than 99.9999 % of DTW computations

Lower bounds for DTW

LB_{KimFL} [Kim et al., 2001]

- ▶ Let $\mathbf{x}$ and $\mathbf{y}$ be two time series of respective lengths M and N
- ▶ Knowing that the first and last points of sequences should be aligned in DTW, i.e. $(i_1, j_1) = (1, 1)$ and $(i_{K_p}, j_{K_p}) = (M, N)$ we know that

$$\begin{aligned} d_{DTW}^2(\mathbf{x}, \mathbf{y}) &= \min_{P \in \mathcal{P}} \sum_{k=1}^{K_p} (x[i_k] - y[j_k])^2 \\ &\geq (x[1] - y[1])^2 + (x[M] - y[N])^2 \end{aligned}$$

$$LB_{KimFL}(\mathbf{x}, \mathbf{y}) = \sqrt{(x[1] - y[1])^2 + (x[M] - y[N])^2}$$

- ▶ This naive lower bound is easy to compute and can provide a first rough estimation

How to use the lower bound property to avoid DTW computations ?




1. Given a pattern $\mathbf{p}$ of length N_p and a time series $\mathbf{x}$ of length $N > N_p$ we first compute $LB_{KimFL}(\mathbf{p}, \mathbf{x}[n : n + N_p - 1])$ for each n (it is very cheap !)
2. Then, we process iteratively.
 - 2.1 **Initialization.** We define the set $\mathcal{N}$ of acceptable samples (initialized as the whole $1 \dots N - N_p + 1$ interval), and we initialize the minimum distance value d^* as $+\infty$
 - 2.2 We select (randomly or uniformly) one or a few samples $n_1, n_2, \dots$ in $\mathcal{N}$ for which we compute **the real DTW values** : $DTW(\mathbf{p}, \mathbf{x}[n : n + N_p - 1])$ and we **update the current minimum distance value** d^*
 - 2.3 All samples n for which

$$LB_{KimFL}(\mathbf{p}, \mathbf{x}[n : n + N_p - 1]) > d^*$$

are removed from $\mathcal{N}$

- 2.4 Reiterate the two previous steps

Results



By computing only one DTW (the first one), 10% of the computations can be avoided (blue = true DTW, green = LB_{KimFL})

Lower bounds for constrained DTW

LB_{Keogh} [Keogh et al., 2004]

- ▶ Let $\mathbf{x}$ and $\mathbf{y}$ be two time series of respective lengths M and N
- ▶ Knowing that we use a constrained DTW with width λ , we know that sample i_k cannot be associated to j_k if $|i_k - j_k| > \lambda$
- ▶ Given a signal $\mathbf{x}$, we can compute the sliding max/min on the window of possible matches

$$u[j] = \max_{j-\lambda \leq i \leq j+\lambda} x[i] \quad l[j] = \min_{j-\lambda \leq i \leq j+\lambda} x[i]$$

- ▶ Then, using the constrained DTW assumption, we have


$$l[j_k] \leq x[i_k] \leq u[j_k]$$

since $j_k - \lambda \leq i_k \leq j_k + \lambda$

- ▶ We therefore have

$$l[j_k] - y[j_k] \leq x[i_k] - y[j_k] \leq u[j_k] - y[j_k]$$

Example



The envelope of $\mathbf{x}$: at each position j , the maximum and the minimum over the window of possible matches

Lower bounds for constrained DTW

LB_{Keogh} [Keogh et al., 2004]

$$l[j_k] - y[j_k] \leq x[i_k] - y[j_k] \leq u[j_k] - y[j_k]$$

- ▶ Then, if $y[j_k] < l[j_k]$ we know that

$$(x[i_k] - y[j_k])^2 \geq (l[j_k] - y[j_k])^2$$

- ▶ Similarly when $y[j_k] > u[j_k]$ we know that

$$(x[i_k] - y[j_k])^2 \geq (u[j_k] - y[j_k])^2$$

- ▶ In other cases, we only know that

$$(x[i_k] - y[j_k])^2 \geq 0$$

- ▶ We have therefore a lower bound for each $(x[i_k] - y[j_k])^2$, that **only depends** on the comparison of $y[j_k]$ to $l[j_k]$ or $u[j_k]$

Lower bounds for constrained DTW

$$LB_{Keogh} \text{ [Keogh et al., 2004]}$$

- ▶ Knowing that

$$d_{DTW}^2(\mathbf{x}, \mathbf{y}) = \min_{p \in \mathcal{P}} \sum_{k=1}^{K_p} (x[i_k] - y[j_k])^2$$

- ▶ We have that


$$d_{DTW}^2(\mathbf{x}, \mathbf{y}) \geq \sum_{\substack{1 \leq j \leq N \\ y[j] < l[j]}} (l[j] - y[j])^2 + \sum_{\substack{1 \leq j \leq N \\ y[j] > u[j]}} (u[j] - y[j])^2$$

- ▶ And finally a new lower bound

$$LB_{Keogh}(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{\substack{1 \leq j \leq N \\ y[j] < l[j]}} (l[j] - y[j])^2 + \sum_{\substack{1 \leq j \leq N \\ y[j] > u[j]}} (u[j] - y[j])^2}$$




- ▶ This lower bound can be used exactly like the previous lower bound

Example



Only the parts of y lying outside the envelope contribute to the bound. The shaded areas are the terms of the sum

Results



($\lambda = 50$). By computing only the 50 first DTW, 80% of the computations can be avoided (blue = true DTW, green = LB_{Keogh})

Contents

1. Problem statement
2. Comparing time series
3. Detecting patterns in time series
4. Learning patterns from time series
 - 4.1 Distance-based pattern extraction
 - 4.2 Dictionary-based pattern extraction
5. Conclusion


Pattern extraction

- ▶ We have seen methods to search for patterns in a time series: these methods need a predefined dictionary of templates
- ▶ In practice, one can be interested in the reverse question: **how can I detect and extract patterns from a time series ?**
- ▶ **Unsupervised task:** no prior knowledge except for the average duration of the researched patterns L
- ▶ **What is a pattern ?** Difficult question : repetitive shapes, notion of periodicity, etc.

Distance-based pattern extraction


- ▶ Intuitively, we can re-use results from the previous part to automatically detect patterns
- ▶ **A pattern is a subsequence that is likely to be found several times in the whole time series**
- ▶ By computing a sliding distance between this subsequence and all subsequences in the time series, it should appear clearly that one (or several) other subsequences are very **close**
- ▶ **Naive solution:** compute all distances by brute force. Can we afford it?

Example



Here, we computed the distance profile of the sequence displayed in red: it is clear that this sequence appears 3 times in the time series. It might be a pattern!

Example



Here, we computed the distance profile of the sequence displayed in red: no obvious structure in the distance profile (except for the exact correspondence): unlikely to be a pattern!


Matrix profile

- ▶ In fact, patterns are detectable by only looking at the **minimal distance with all subsequences**
- ▶ Beware of **trivial matches** ! Only compare with subsequences that do not overlap with the subsequence of interest
- ▶ **Matrix profile** [Yeh et al., 2016] : given a pattern length L , compute

$$m[n] = \min_{i > n+L \text{ OR } i < n-L} d(x[n : n + L - 1], x[i : i + L - 1])$$

- ▶ Small matrix profile values indicate that the subsequence has been found elsewhere in the time series, suggesting that it could be a pattern
- ▶ Efficient computation with techniques already mentioned (fast sliding distance computation)

Example



In red : subsequence with minimal matrix profile value. In green : closest subsequence (according to the normalized Euclidean distance)

From matrix profile to motif discovery

- ▶ The matrix profile identifies a **pair**: the two most similar non-overlapping subsequences (motif pair problem), but applications require **all** occurrences of a recurring pattern
- ▶ **Motif set problem**: identify sets of subsequences that represent the same pattern
- ▶ This problem can be solved using a procedure that identifies all sufficiently close subsequences (by applying a threshold to the distance profile) and then repeats the process until the desired number of motif sets is found.
- ▶ These parameters (threshold + number of motif sets) can have a significant impact on the detection, and they are difficult to adjust. See mini-projects !



Dictionary-based pattern extraction

- ▶ Another approach consists in formulating this problem as a **dictionary learning** optimization problem
- ▶ Given an input time series $\mathbf{x}$, learn a **dictionary of K patterns $\mathbf{d}_k$** of length L
- ▶ These patterns can be activated : **activations $\mathbf{z}_k$** of length $N - L + 1$

$z_k[n] \neq 0$ if pattern $\mathbf{d}_k$ is activated at time n

[Grosse et al., 2007 ; Wohlberg, 2014]

Convolutional dictionary learning



Convolutional dictionary learning

Given a time series $\mathbf{x}$, number of patterns K and pattern length L , learn

- ▶ Patterns $\mathbf{d}_k$ of length L
- ▶ Activation signals $\mathbf{z}_k$ of length $N - L + 1$

$$x[n] = \sum_{k=1}^K (\mathbf{z}_k * \mathbf{d}_k)[n] + e[n]$$

Optimization problem

$$\min_{\substack{(\mathbf{d}_k), (\mathbf{z}_k) \\ \forall k, \|\mathbf{d}_k\|_2^2 \leq 1}} \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2 + \lambda \sum_{k=1}^K \|\mathbf{z}_k\|_1$$

- ▶ **Normalization constraint** for the dictionary atoms $\mathbf{d}_k$, that prevents numerical instabilities (otherwise setting $\alpha \mathbf{d}_k$ and $\alpha^{-1} \mathbf{z}_k$ gives the same result)
- ▶ **Sparsity constraint** for the activations $\mathbf{z}_k$, that improves the interpretability of the learned patterns (see Lecture 3)

Not convex with respect to the couple $(\mathbf{d}_k), (\mathbf{z}_k)$ but convex when the subproblems are taken individually

Alternated resolution

Alternated resolution of two subproblems

Dictionary learning

$$\mathbf{D}^* = \underset{\substack{\mathbf{D}=(\mathbf{d}_1,\dots,\mathbf{d}_K) \\ \forall k, \|\mathbf{d}_k\|_2^2 \leq 1}}{\operatorname{argmin}} \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2$$

Convolutional sparse coding

$$\mathbf{Z}^* = \underset{\mathbf{Z}=(\mathbf{z}_1,\dots,\mathbf{z}_K)}{\operatorname{argmin}} \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2 + \lambda \sum_{k=1}^K \|\mathbf{z}_k\|_1$$

Alternated resolution

$$\min_{\substack{(\mathbf{d}_k), (\mathbf{z}_k) \\ \forall k, \|\mathbf{d}_k\|_2 \leq 1}} \underbrace{\left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2}_{f(\mathbf{Z}, \mathbf{D})} + \lambda \sum_{k=1}^K \|\mathbf{z}_k\|_1$$

- ▶ Both these problems can be solved with **Proximal Gradient Descent** algorithms
- ▶ **Two main steps** :
 1. **Gradient descent step** w.r.t. $\nabla_{\mathbf{D}} f(\mathbf{Z}, \mathbf{D})$ or $\nabla_{\mathbf{Z}} f(\mathbf{Z}, \mathbf{D})$
 2. **Proximal step** to **project** the update on the constraint set
- ▶ The main question is therefore: **how can we easily compute the gradients?**

Formulation of the gradients

$$f(\mathbf{Z}, \mathbf{D}) = \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2$$

- By using the **convolution theorem** and the **Parseval theorem**, we have that

$$f(\mathbf{Z}, \mathbf{D}) \propto \left\| \hat{\mathbf{x}} - \sum_{k=1}^K \hat{\mathbf{z}}_k \odot \hat{\mathbf{d}}_k \right\|_2^2$$

where $\hat{\cdot}$ denotes the Discrete Fourier Transform (DFT) and $\odot$ the component-wise product

- As $\mathbf{u} \odot \mathbf{v} = \text{diag}(\mathbf{u})\mathbf{v}$, this expression can be rewritten in the matrix form as

$$f(\mathbf{Z}, \mathbf{D}) = \left\| \hat{\mathbf{x}} - \hat{\mathbf{D}}\hat{\mathbf{z}} \right\|_2^2$$

- The gradient of this quantity is easy to compute w.r.t. $\hat{\mathbf{D}}$ and $\hat{\mathbf{z}}$
- Updates for $\mathbf{D}$ and $\mathbf{Z}$ can then be estimated by performing inverse DFT

Proximal operators

- ▶ For the unit ball constraint (dictionary), the proximal operator writes as

$$\text{proj}_{\|\cdot\|_2 \leq 1}(\mathbf{y}) = \frac{\mathbf{y}}{\max(\|\mathbf{y}\|_2, 1)}$$

Projection on the unit ball

- ▶ For the L1-sparsity constraint (activations), the proximal operator writes as

$$\mathcal{S}_\gamma(\mathbf{y})[n] = \text{sign}(y[n]) \times \max(|y[n]| - \gamma, 0)$$

Soft thresholding operator

Dictionary learning

$$\mathbf{D}^* = \underset{\substack{\mathbf{D}=(\mathbf{d}_1, \dots, \mathbf{d}_K) \\ \forall k, \|\mathbf{d}_k\|_2 \leq 1}}{\operatorname{argmin}} \underbrace{\left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2}_{f(\mathbf{Z}, \mathbf{D})}$$

► **Basic approach : Proximal Gradient Descent with fixed $\mathbf{Z}$**

1. Gradient step :

$$\mathbf{D} \leftarrow \mathbf{D} - \alpha \nabla_{\mathbf{D}} f(\mathbf{Z}, \mathbf{D})$$

2. Proximal projection step :

$$\mathbf{d}_k \leftarrow \operatorname{proj}_{\|\cdot\|_2 \leq 1}(\mathbf{d}_k) = \frac{\mathbf{d}_k}{\max(\|\mathbf{d}_k\|_2, 1)}$$

► **Other approaches** : Alternating Direction Method of Multipliers (ADMM), K-SVD... (see [Mairal et al., 2010] and Lecture 3 for more details)

Convolutional sparse coding

$$\mathbf{Z}^* = \underset{\mathbf{Z}=(\mathbf{z}_1, \dots, \mathbf{z}_K)}{\operatorname{argmin}} \underbrace{\left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2}_{f(\mathbf{Z}, \mathbf{D})} + \lambda \underbrace{\sum_{k=1}^K \|\mathbf{z}_k\|_1}_{\psi(\mathbf{Z})}$$

- **Basic approach : Iterative Soft Thresholding Algorithm (ISTA)** with fixed $\mathbf{D}$ [Daubechies et al., 2004]

1. **Gradient step**


$$\mathbf{Z} \leftarrow \mathbf{Z} - \alpha \nabla_{\mathbf{Z}} f(\mathbf{Z}, \mathbf{D})$$

2. **Proximal step**


$$\mathbf{Z} \leftarrow \mathcal{S}_{\lambda\alpha}(\mathbf{Z})$$

- **Other approaches:** Alternating Direction Method of Multipliers (ADMM), Fast Iterative Soft Thresholding Algorithm (FISTA), Coordinate Descent (CD) (see [Mairal et al., 2010] and Lecture 3 for more details)


Results




Results




Results




Results




Results




Results



Results




Results




One pattern

Results



Two patterns

Results



Mixture of patterns

Summary

- ▶ Contrary to distance-based pattern extraction techniques, **CDL** can detect **mixture of patterns** thanks to the additive model
- ▶ CDL is traditionally quite **sensitive to initialization and/or hyperparameters** : random initializations (normalized gaussian noise), randomly chosen data frames, Lipschitz constant for gradient descent...
- ▶ CDL can also be used to search patterns from a fixed dictionary of templates : in this case, only perform the convolutional sparse coding step

Contents

1. Problem statement
2. Comparing time series
3. Detecting patterns in time series
4. Learning patterns from time series
5. Conclusion

Conclusion

▶ Comparing time series and distances

- ▶ **DTW** remains the reference and the most popular elastic distance
- ▶ In the context of clustering, where metric properties are important (particularly for calculating the barycenter), certain **measures based on edit distances (MSM, TWE)** may prove to be more accurate (see a recent review)
- ▶ Differentiable relaxations allow elastic distances to be used as **training losses**

▶ Pattern detection

- ▶ Exact search under (possibly normalized) Euclidean distance in $\mathcal{O}(N \log N)$
- ▶ Under DTW: lower bounds prune $> 99.99\%$ of the computations, scales to trillions of subsequences

▶ Pattern extraction

- ▶ The **Matrix profile** solved the **motif pair** problem
- ▶ There is no single method that stands out for solving the **motif set** problem in the general case (multiple patterns, whose number, length, and cardinality are unknown).
- ▶ Recent methods try to reduce the number of parameters (e.g. by using cardinality-based or topological approaches) and the evaluation is tricky

▶ Dictionary learning

- ▶ Proximal methods, FFT-based gradients: optimization
- ▶ Recent progress on **guarantees** (ℓ_0 exact solutions) and **robustness**

Benchmarks

TSMD: benchmark on univariate motif discovery [Guerrini et al., 2025]: 11 algorithms, 8 annotated real datasets (569 time series) and a synthetic generator designed to isolate individual challenges

Conclusions

- ▶ **No method dominates:** a few stand out slightly on real data, but performance varies a lot across datasets
- ▶ Robustness properties **differ considerably:** depending of the criterion of evaluation (sensibility to a large number of motifs, to noise, to intra-motif length variation, to trends, and scalability...) the choice of which method to use varies
- ▶ Performance **collapses** when motifs occupy a small fraction of the series : precisely the realistic industrial setting
- ▶ **Huge challenge:** extension to the multivariate case

A few reference papers

- ▶ Guerrini, V., Germain, T., Truong, C., Oudre, L., & Boniol, P. (2025). Time series motif discovery: a comprehensive evaluation. *Proceedings of the VLDB Endowment (PVLDB)*, 18(7), 2226-2239. : **benchmark for motif discovery**
- ▶ Holder, C., Middlehurst, M., & Bagnall, A. (2024). A review and evaluation of elastic distance functions for time series clustering. *Knowledge and Information Systems*, 66(2), 765-809. : **benchmark for clustering**
- ▶ Middlehurst, M., Schäfer, P., & Bagnall, A. (2024). Bake off redux: a review and experimental evaluation of recent time series classification algorithms: M. Middlehurst et al. *Data Mining and Knowledge Discovery*, 38(4), 1958-2031. : **benchmark for classification**
- ▶ Paparrizos, J., Li, H., Yang, F., Wu, K., d'Hondt, J. E., & Papapetrou, O. (2024). A survey on time-series distance measures. *arXiv preprint arXiv:2412.20574*. : **survey on time series distances**
- ▶ Shifaz, A., Pelletier, C., Petitjean, F., & Webb, G. I. (2023). Elastic similarity and distance measures for multivariate time series. *Knowledge and Information Systems*, 65(6), 2665-2698. : **survey on multivariate distances**

Code and resources

- ▶ `aeon` - <https://www.aeon-toolkit.org>
Most elastic distances (DTW, edit distances...)
- ▶ `tslearn` - <https://tslearn.readthedocs.io>
DTW, soft-DTW, barycenters (DBA, soft-DBA)
- ▶ `stumpy` - <https://stumpy.readthedocs.io>
Matrix profiles and variants (STOMP, GPU, streaming, multidimensional...)
- ▶ `TSMD` - <https://github.com/grrvrlr/TSMD>
Benchmark for motif discovery (k-Motiflets, GrammarViz, PEPA...) with data
- ▶ `alphacsc` - <https://alphacsc.github.io>
Convolutional sparse coding

Datasets

Classification / similarity search

- ▶ **UCR/UEA archive** - <https://www.timeseriesclassification.com>
128 univariate + 30 multivariate datasets

Motif discovery

- ▶ **TSMD** - <https://github.com/grrvlr/TSMD>
8 datasets, 569 times series with annotated motifs and a synthetic generator

Physiological signals

- ▶ **PhysioNet** - <https://physionet.org>
MIT-BIH arrhythmia database, PPG, EEG, and many others

References

- ▶ Berndt, D. J., & Clifford, J. (1994, July). Using dynamic time warping to find patterns in time series. In KDD workshop (Vol. 10, No. 16, pp. 359-370).
- ▶ Sakoe, H., & Chiba, S. (1978). Dynamic programming algorithm optimization for spoken word recognition. *IEEE transactions on acoustics, speech, and signal processing*, 26(1), 43-49.
- ▶ Itakura, F. (1975). Minimum prediction residual principle applied to speech recognition. *IEEE Transactions on acoustics, speech, and signal processing*, 23(1), 67-72.
- ▶ Mueen, A., Viswanathan, K., Gupta, C. K., & Keogh, E. (2015). The fastest similarity search algorithm for time series subsequences under Euclidean distance. <http://www.cs.unm.edu/~mueen/FastestSimilaritySearch.html>
- ▶ Rakthanmanon, T., Campana, B., Mueen, A., Batista, G., Westover, B., Zhu, Q., ... & Keogh, E. (2012, August). Searching and mining trillions of time series subsequences under dynamic time warping. In *Proceedings of the 18th ACM SIGKDD international conference on Knowledge discovery and data mining* (pp. 262-270).
- ▶ Kim, S. W., Park, S., & Chu, W. W. (2001, April). An index-based approach for similarity search supporting time warping in large sequence databases. In *Proceedings 17th International Conference on Data Engineering* (pp. 607-614). IEEE.
- ▶ Keogh, E., & Ratanamahatana, C. A. (2004). Exact indexing of dynamic time warping. *Knowledge and information systems*, 7(3), 358-386.
- ▶ Yeh, C. C. M., Zhu, Y., Ulanova, L., Begum, N., Ding, Y., Dau, H. A., ... & Keogh, E. (2016, December). Matrix profile I: all pairs similarity joins for time series: a unifying view that includes motifs, discords and shapelets. In *2016 IEEE 16th international conference on data mining (ICDM)* (pp. 1317-1322). Ieee.
- ▶ Grosse, R., Raina, R., Kwong, H., & Ng, A. Y. (2007, July). Shift-invariant sparse coding for audio classification. In *Proceedings of the Twenty-Third Conference on Uncertainty in Artificial Intelligence* (pp. 149-158).
- ▶ Wohlberg, B. (2014, May). Efficient convolutional sparse coding. In *2014 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (pp. 7173-7177). IEEE.
- ▶ Mairal, J., Bach, F., Ponce, J., & Sapiro, G. (2010). Online learning for matrix factorization and sparse coding. *Journal of Machine Learning Research*, 11(1).
- ▶ Daubechies, I., Defrise, M., & De Mol, C. (2004). An iterative thresholding algorithm for linear inverse problems with a sparsity constraint. *Communications on Pure and Applied Mathematics: A Journal Issued by the Courant Institute of Mathematical Sciences*, 57(11), 1413-1457.

Lecture 1: list of possible topics/projects (1/2)

Elastic distances & differentiable losses

- Project 1.1** Cuturi, M., & Blondel, M. (2017, July). Soft-dtw: a differentiable loss function for time-series. In International conference on machine learning (pp. 894-903). PMLR.
How to make DTW differentiable using optimal transport
- Project 1.2** Le Guen, V., & Thome, N. (2019). Shape and time distortion loss for training deep time series forecasting models. Advances in neural information processing systems, 32.
How to construct a DTW-based loss for deep learning
- Project 1.3** Herrmann, M., & Webb, G. I. (2023). Amercing: An intuitive and effective constraint for dynamic time warping. Pattern Recognition, 137, 109333.
How to constrain DTW more effectively than with the Sakoe-Chiba band
- Project 1.4** Holder, C. & Bagnall, A. (2026). Soft-MSM: Differentiable Context-Aware Elastic Alignment for Time Series, arXiv:2605.00069.
How to make MSM (another elastic distance) differentiable for use as a deep learning loss

Shapelets

- Project 1.5** Hills, J., Lines, J., Baranauskas, E., Mapp, J., & Bagnall, A. (2014). Classification of time series by shapelet transformation. Data Mining and Knowledge Discovery, 28(4), 851-881.
How to use patterns for time series classification
- Project 1.6** Guillaume, A., Vrain, C., & Elloumi, W. (2022, June). Random dilated shapelet transform: A new approach for time series shapelets. In International Conference on Pattern Recognition and Artificial Intelligence (pp. 653-664)
How to combine shapelets and random convolutional kernels using dilation

Matrix Profile and variants

- Project 1.7** Linardi, M., Zhu, Y., Palpanas, T., & Keogh, E. (2020). Matrix profile goes MAD: variable-length motif and discord discovery in data series. Data Mining and Knowledge Discovery.
How to extend the matrix profile approach to variable lengths motifs
- Project 1.8** Yeh, C. C. M., Kavantzaz, N., & Keogh, E. (2017, November). Matrix profile vi: meaningful multidimensional motif discovery. In 2017 IEEE international conference on data mining (ICDM) (pp. 565-574). IEEE.
How to extend the matrix profile approach to multivariate time series
- Project 1.9** Alaae, S., Mercer, R., Kamgar, K., & Keogh, E. (2021). Time series motifs discovery under DTW allows more robust discovery of conserved structure. Data Mining and Knowledge Discovery, 35(3), 863-910.
How to combine the Matrix Profile with DTW for more robust motif discovery

Lecture 1: list of possible topics/projects (2/2)

Motif discovery

- Project 1.10** Li, Y., & Lin, J. (2010). Approximate variable-length time series motif discovery using grammar inference. In Proceedings of the Tenth International Workshop on Multimedia Data Mining (pp. 1-9).
How to discover motifs using grammar inference on a discretized series
- Project 1.11** Schäfer, P., & Leser, U. (2022). Motiflets: Simple and accurate detection of motifs in time series. Proceedings of the VLDB Endowment, 16(4), 725-737.
How to find motifs without setting a similarity threshold
- Project 1.12** Germain, T., Truong, C., & Oudre, L. (2024). Persistence-based motif discovery in time series. IEEE Transactions on Knowledge and Data Engineering, 36(11), 6814-6827.
How to use topological data analysis to discover motifs

Convolutional Dictionary Learning

- Project 1.13** Pilastre, B., Silva, G., Boussouf, L., d'Escrivan, S., Rodriguez, P., & Tournet, J. Y. (2020, May). Anomaly Detection in Mixed Time-Series Using A Convolutional Sparse Representation With Application To Spacecraft Health Monitoring. In ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP) (pp. 3242-3246). IEEE.
How to use convolutional dictionary learning for detecting anomalies
- Project 1.14** La Tour, T. D., Moreau, T., Jas, M., & Gramfort, A. (2018). Multivariate convolutional sparse coding for electromagnetic brain signals. In Advances in Neural Information Processing Systems (pp. 3292-3302).
How to use convolutional dictionary learning to study the brain
- Project 1.15** Truong, C., & Moreau, T. (2024). Convolutional sparse coding for time series via a ℓ_0 penalty: an efficient algorithm with statistical guarantees. Statistical Analysis and Data Mining, 17(6), e70000.
How to obtain exact CDL solutions with statistical guarantees
- Project 1.16** Yehya, J., Benbakoura, M., Allain, C., Malezieux, B., Kowalski, M., & Moreau, T. (2025). RoseCDL: Robust and scalable convolutional dictionary learning for rare event and anomaly detection. arXiv preprint arXiv:2509.07523.
How to make CDL robust and scalable for rare-event detection in long signals